

Masked LARK

Masked Learning, Aggregation & Reporting worKflow

Proposal for Decentralized Aggregation for Conversion Reporting & Modeling

Denis [Charles](#), Joel Pfeiffer, Mehul Parsana, Erik Anderson - Microsoft

Goals of the talk

- Introduce technical aspects of [Masked LARK](#)
 - Aim of the proposal.
 - Algorithmic & Technical details of Masked LARK.
 - How it differs from existing proposals.
 - Known limitations.
- We will not cover API/Workflow details here.
 - [ToDo in a future session.](#)

Background - 3rd party cookies deprecation

Removal of 3rd party cookies impacts the following:

- Fraud detection
- Targeting
- Conversion reporting and modeling

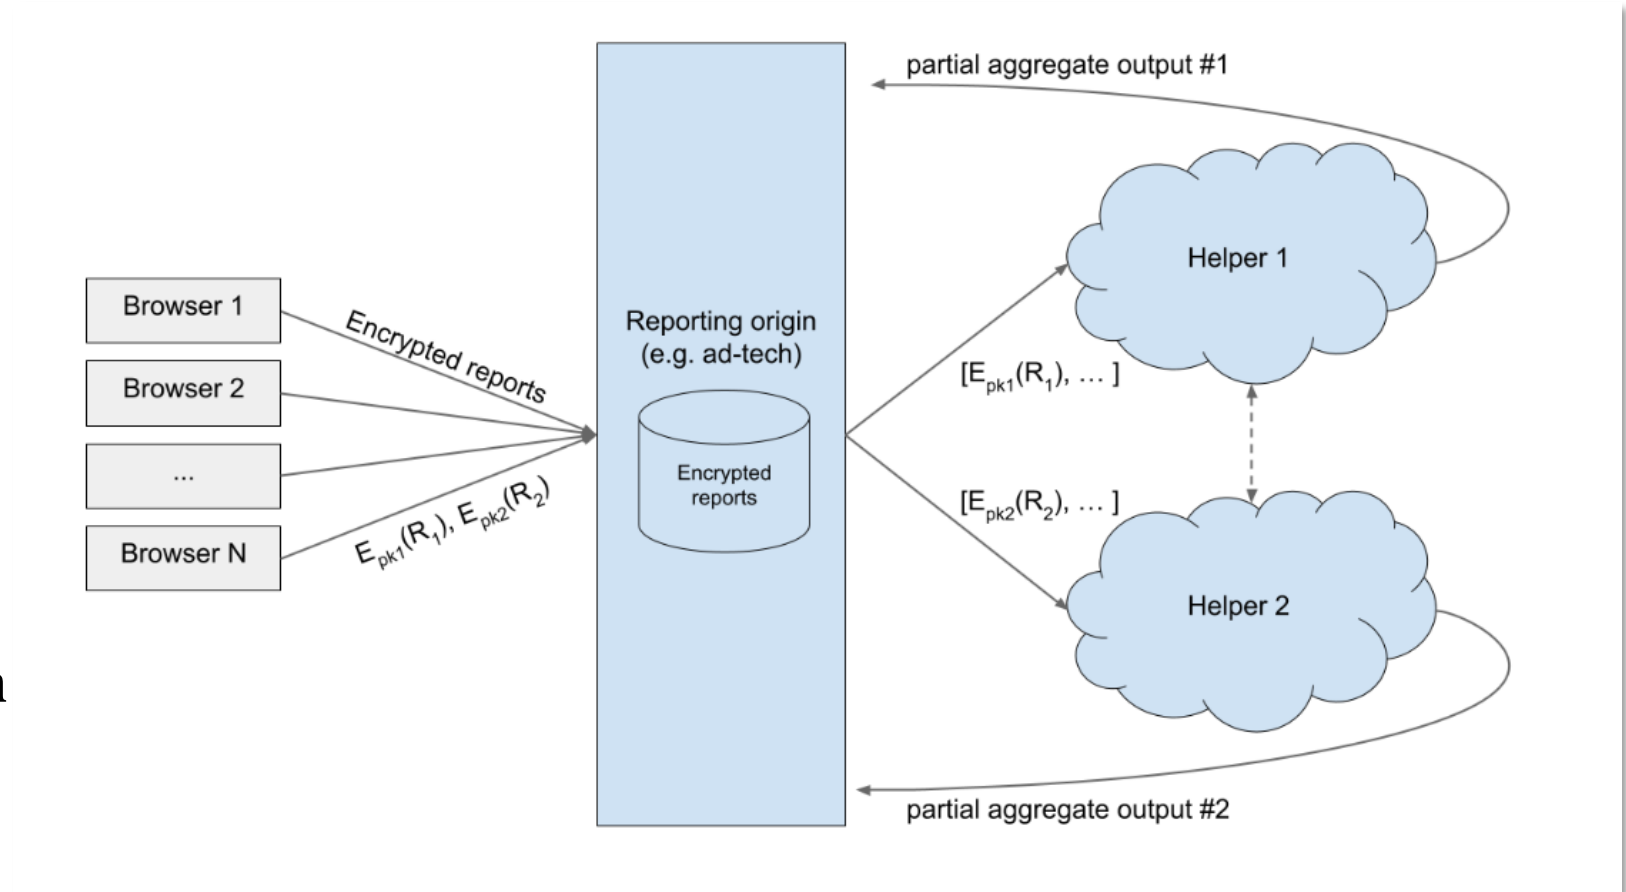
Some Proposals:

- Fraud – [Trust tokens](#)
- For targeting
 - [FLoC](#)
 - [Turtledove](#)
 - [Parakeet](#)
- Conversion reporting – [Multi-Browser Aggregation Service](#)

Google Proposal for Conversion reporting

Idea:

- Implement trusted mediator abstraction with multiple semi-trusted helpers.
- Use secure Multi-party computation and differential privacy to aggregate data.
- In particular, **no single** helper has data about final aggregates.



Google Proposal for Conversion reporting

Pros:

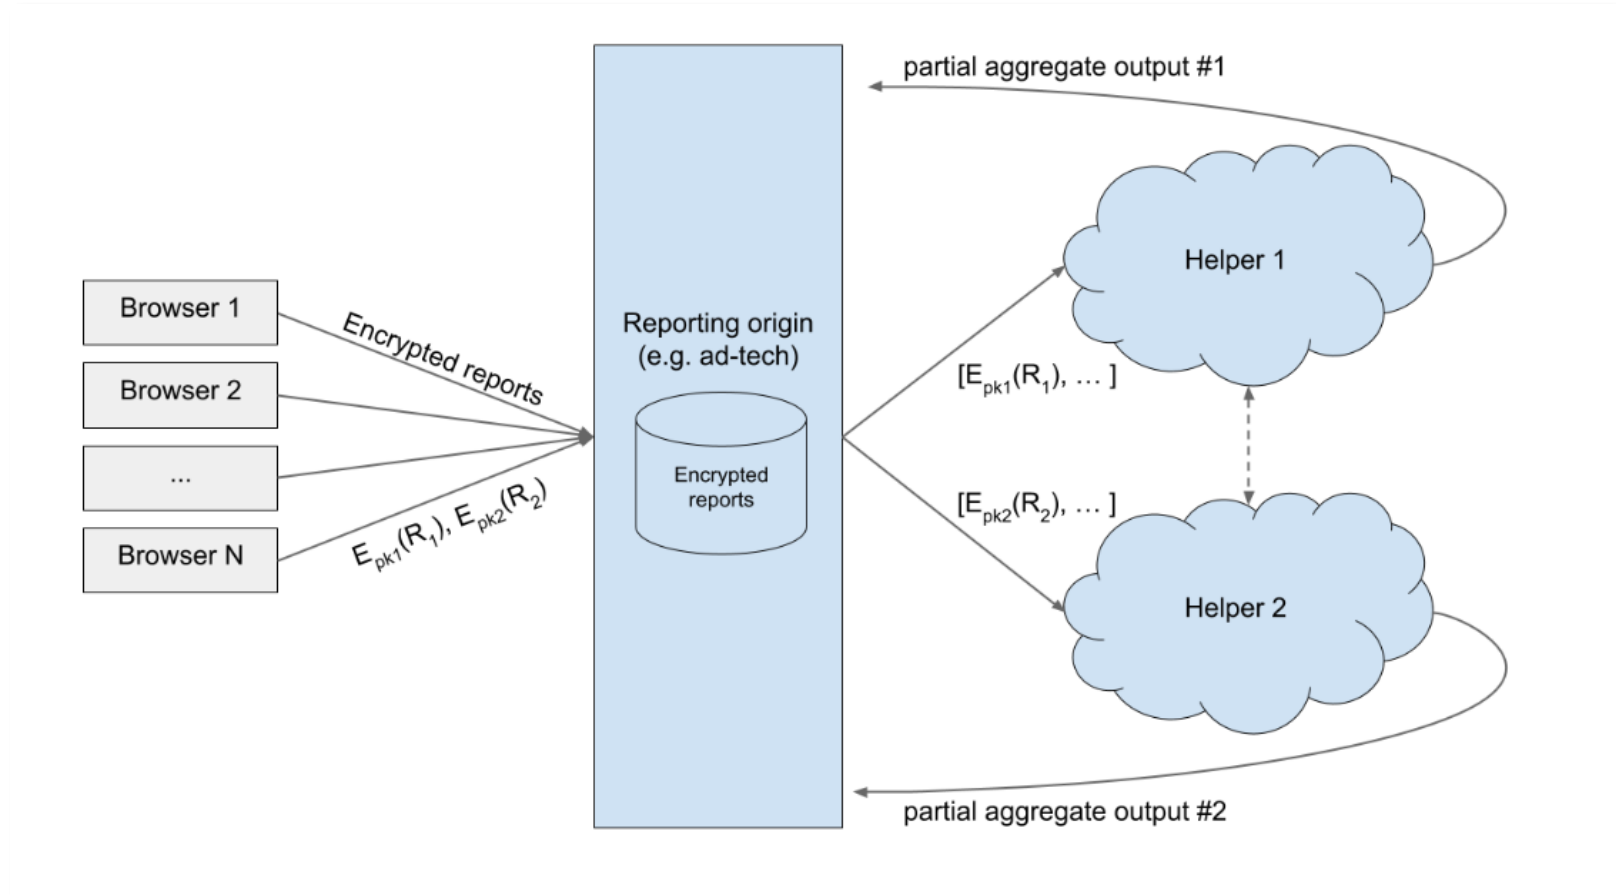
- Segregated helpers implementing aggregation more palatable than single trusted mediator.
- MPC implementation is easier with always available trusted parties.

Limitations:

- Handles aggregation for reporting needs but does not* address modeling.

Our Goal:

- Build on top of proposal to address modeling needs.
- Focus on simple ideas for privacy (minimal crypto – only secret sharing, use only finite rings).



* Can handle with FL & secure aggregation (has disadvantages, discuss later).

General View: Differentially private Map-Reduce

- One can view the aggregation service as implementing a differentially private Map-Reduce framework with semi-trusted helpers.
- Browsers apply a secure “Map” operation.
 - They take local user activity and create keys and data for aggregation.
- Helpers apply a differentially private “Reduce” operation.
 - Any secure MPC function that is differentially private can be a reduce operation. Some examples are sum, approximate rank in sorted list etc.
 - Clarifying requirements for helpers allow us to onboard more helpers.
- Ad Network only sees output of the full process.
- Honesty of helpers can be tested by independent parties (including the ad network).
- Formalization of framework allows users to reason about the system.

Requirements for Helpers

Requirements: Very similar to secure aggregation (reduce operation) with additional differential privacy requirements.

- In some time period T
- Each user (browser) has $\mathbf{v}_i \in \mathbb{Z}^d \subseteq \mathbb{R}^d$, this can be client level statistics on user activity
- Helpers compute and output $\mathbf{s} \approx \sum_i \mathbf{v}_i \equiv \mathbf{w}$ with the following properties:
 - **[Noisy aggregation.]** $\mathbf{s}_i \in_R D(\mu_i = \mathbf{w}_i, \sigma_i = \epsilon \mu_i)$.
 - **[Sparse reporting constraint/ k -Anonymity.]** Another useful property is that if there are $< k$ clients reporting with $\mathbf{v}_{i\ell} \neq 0$, then $\mathbf{s}_\ell = \emptyset$.
 - In other words, if less than k clients share an aggregation key that key is garbage collected.
- As part of the protocol, AdNetwork publicly discloses ϵ and k .

Notes:

- One difference from Google proposal is there is no explicit hiding of the aggregation keys.
- Technical point, σ_i need to be specified per independent observation $\hat{\sigma}_i$. Helpers will sample from $D(0, \sigma_i = \hat{\sigma}_i \sqrt{k})$.

Example: Counting conversions

- Suppose we wish to compute *Query_X_Advertiser* level #Clicks & #Conversions, with sparse reporting constraint of N users.
- Assume we have 2 helpers.
- In Time window T, browser i creates key $k = Enc_{AdNet}(q, a)$ and sets value $\tilde{v} = \langle c_i, conv_i \rangle$.
 - If we have two helpers and suppose we work with ring $\mathbb{Z}/m\mathbb{Z}$, browser generates random values $\langle r_1, r_2 \rangle$ and creates $v_{i,1} = \langle r_1 + c_i, r_2 + conv_i \rangle \bmod m$ and $v_{i,2} = \langle -r_1, -r_2 \rangle \bmod m$.
 - Browser encrypts $v_{i,1}$ and $v_{i,2}$ using public key of the helpers.
- Aggregation service computes $\langle \sum clicks_i, \sum conv_i \rangle + \epsilon_{noise} \approx \sum_{i,j} v_{i,j}$.
- If number of *distinct users* reporting information for key, $k = (q, a)$, is $\leq N$ we get #clicks = \emptyset and #conv = \emptyset .
 - **Note:** Since each helper sees data from every user reporting a key. They can individually apply the sparse reporting constraint.
- This encourages aggregations to not be granular. Perhaps use query categories for aggregation or require larger time window.

Model training

- For differentiable models M :
 - For Model training set $v_i = \partial \text{loss}(f_i, \text{label}_i, M)$, f_i = local feature vector.
 - If we use models with count features, we can update aggregates by the protocol directly.
- **Potential approaches:**
 - Since v_i in this case depends on actual label and model, it needs to be computed on user's device. This implies a **federated learning** setup is required.
 - May be expensive to do on user's device as number of models can be large.
 - If user trusts helper with label, gradient computation can be done at helper.
 - Trust in single helper is too strong of an assumption.
 - Another alternative is for user to send true label to aggregate with probability p and random label with probability $1 - p$.
 - Not great from privacy standpoint.

Our solution - Masked aggregation

Outline of solution in simplified setting.

For binary labels (like conversion models), we can do the following.

- Browser sends both $\langle f_i, \text{label} = 0 \rangle$ and $\langle f_i, \text{label} = 1 \rangle$ to each helper.
 - **Note:** f_i is known to AdNetwork, so AdNetwork can provide feature vector to helper. This saves communication cost on user side.
- Helpers compute two gradient vectors $\mathbf{g}_{i,0} = \partial \text{loss}(f_i, \mathbf{0}, \mathbf{M})$ and $\mathbf{g}_{i,1} = \partial \text{loss}(f_i, \mathbf{1}, \mathbf{M})$, where we encode these gradients in $(\mathbb{Z}/m\mathbb{Z})^d$.
- Separately, browser generates $\langle \alpha_{i,0}, \alpha_{i,1} \rangle$ and $\langle \beta_{i,0}, \beta_{i,1} \rangle$.
 - If 0 is the true label, $\alpha_{i,0} + \alpha_{i,1} = 1$ and 0 otherwise.
 - Similarly, if 1 is the true label $\beta_{i,0} + \beta_{i,1} = 1$ and 0 otherwise.
 - Here $\alpha_{i,k} = r_i$ where r_i is random in $\mathbb{Z}/m\mathbb{Z}$, and we set $\alpha_{i,0} + \alpha_{i,1} = \mathbf{1}[\text{label} = 0] \pmod{m}$ as above. Similarly, we set $\beta_{i,k}$.
- Browser sends $\alpha_{i,k}$ and $\beta_{i,k}$ to Helper k .
- Finally, helpers compute gradient $\mathbf{G} = \sum_{i,k} \alpha_{i,k} \mathbf{g}_{i,0} + \beta_{i,k} \mathbf{g}_{i,1}$.
 - For each i , we have by construction that $\sum_k \alpha_{i,k} \mathbf{g}_{i,0} + \beta_{i,k} \mathbf{g}_{i,1} \equiv \text{True gradient}$.
- \mathbf{G} is revealed to AdNetwork **after** the differential privacy constraints are applied. This is used by AdNetwork to update model.

Idea generalizes to non-binary case by introducing “real” gradients and “fake” gradients for aggregation.

Also generalizes to “real” and “fake” features / models, which can protect sensitive ad network data.

Masked aggregation – Abstract setting

Lemma: There is a simple secret sharing protocol to compute $\langle v, w \rangle$ for $v, w \in V$, an R -module (equipped with a bilinear form $\langle \cdot, \cdot \rangle$), with three parties.

Proof: Suppose we have three parties H_1, H_2 and C , and we assume that w belongs to C , and H_1, H_2 both have v .

Now to compute $\langle v, w \rangle$, C does a secret sharing protocol with H_i and shares α_i to H_1 and β_i to H_2 with $\alpha_i \in_R R$: $\alpha_i + \beta_i = w_i$.

The H_i can now compute $\sum \alpha_i v_i, \sum \beta_i v_i$ and the sum is $\langle v, w \rangle$ by bilinearity of the inner product. ■

Corollary: There is a simple secret sharing protocol to compute $\sum_{i \in S} v_i$ where S is secret to one party and v_i are shared with two other parties.

Proof: Apply lemma with $w = \chi(S)$. ■

Masked LARK - Notes

Known issues and potential solutions:

- **Differentially private map-reduce:**
 - Aggregation keys are not hidden but are opaque (encrypted). Helpers can learn if a key was involved in aggregation.
 - Users can insert fake self-cancelling records for aggregation.
 - k -anonymity can be attacked by ad network via ballot stuffing.
 - Ballot stuffing needs ad network to know/guess rare keys.
 - DP on top of aggregation gives another layer of protection.
- **Model training:**
 - Feature vector is not hidden to helper.
 - Work only with dense feature vectors.
 - Users can perturb features or ad network can help by another layer of splitting of feature vectors (communication cost between helpers to sync after first layer update).
 - Label space can leak information in non-binary case.
 - Quantized-label space with randomized rounding.
 - Pollution attacks from users/browsers.
 - Solutions need more expensive crypto.
 - Timing attacks:
 - If users send model updates when activity occurs, this can leak information.
 - Model training need not be synchronous with actual conversion/non-conversion events. Hence real/fake gradients can always be reported at set times for each reporting client.
 - Solution limited to models that are continuous functions trained using SGD.
 - More general models can be trained with general secure MPC. Comes at high cost, several orders of magnitude greater than masked aggregation.